

Engineering of Advanced Software Solutions (EASS)

HIT, Israel

Yossi Eliaz

2022

- Admin
- Technical debt
- Business logic
- Bash and commandline (based on MIT's missing semester)
- Git, GitHub
- Interactive class



- Discord
- Github account
- HW on Github
- Creating a Canvas account
(<https://canvas.instructure.com/>)
- Accepted invitation from AWS Academy and GitHub (I have sent links)
- AWS on Cavnvas
- LinkedIn
- Commandline (WSL)
- Docker
- Moodle (minimal interaction over there)
- Volunteer to summarize the lectures
- Stackoverflow
- Engagment on Discord
- Hackernews

References:

Missing Semester MIT
AWS cloud certificate

What is Technical Debt?

- “In software development, there is always a constant need to balance speed and quality. Some quality will always have to be sacrificed to release features within a reasonable timeframe, so any of these shortcuts will often be tasked as future projects. Those unattended tasks become what is called technical debt.”
- “There are several reasons why technical debt happens. Product owners may focus more on the need to implement and release new features and less on fixing past problems or create a generic enough infrastructure to support future developments. In some severe cases, product owners completely underestimate the outcomes of dealing with poor infrastructure, bugs and poorly designed software.”
- “Ultimately, technical debt can sometimes lead to software users having bad experiences and thereby increasing user churn rates. Together, a lack of developer awareness and task ownership can lead to more technical debt.”

References:

<https://logz.io/blog/technical-debt/>

Business logic

- “**Business rules** are what your non-software developers tell you what your software needs to do.”
- “**Business logic** is the part of your code that specifically implements business rules.”



References:

<https://softwareengineering.stackexchange.com/questions/234251/what-really-is-the-business-logic>

<http://www.ritholtz.com>

- Vim
- Bash
- Git
- Docker

Important tools and commands

- echo, while, find, vars, printenv, htop, shebang, wild cards
- cp, touch, mkdir, ls, uniq, awk, rm
- man man
- brew
- wget
- curl

References:

<https://missing.csail.mit.edu/2020/shell-tools/>

- `docker run`
- `docker ps`
- `docker run -ti`

Must have:

- 1 Problem statement
- 2 Sample code and data
- 3 Spelling, grammar and formatting

Example:

<https://stackoverflow.com/questions/11227809/why-is-processing-a-sorted-array-faster-than-processing-an-unsorted-array>

References:

<https://codeblog.jonskeet.uk/2010/08/29/writing-the-perfect-question/>
<https://stackoverflow.com/help/how-to-ask>

- Checkout github classroom and the first task about git and github
<https://classroom.github.com/classrooms/99552739-eass-hit-2022-part-a>

- S3, EC2, RDS, and EBS modules
- must get 100 on all 4 modules
- grading will be 25% per module

A bit more about EC2 instances and types of hardwares (HW)

EC2 provides secure, resizable compute cloud services. It makes web-scale cloud computing easier and offers HW such as:

- ARM vs. Intel vs. AMD (x86, x86_64)
- GPUs (Nvidia, Intel)
- TPUs (on Google Computing Platform)
- Metal instances on AWS
- FPGA-based nodes

Instance Types and prices (useful links)

<https://aws.amazon.com/ec2/instance-types/>
<https://instances.vantage.sh/>

Instance Types (summary)

- 1 General Purpose
- 2 Compute Optimised
- 3 Memory Optimised
- 4 Accelerated Computing (P instances are for general-purpose GPU applications)

Pricing

There are four ways to pay for EC2 instances: On-Demand, Reserved Instances, and Spot Instances & Per-Second Billing. You can also pay for Dedicated Hosts which provide you with EC2 instance capacity on physical servers dedicated for your use.

Checkout github classroom and the first task about git and github

<https://classroom.github.com/classrooms/99552739-eass-hit-2022-part-a>

How to test our code/system

General approaches for testing

- Static vs. Dynamic
- Passive testing
- White-box vs. Black-box testing

Types of testing coverage metric

- API testing – testing all public and private APIs
- Code coverage – creating tests to satisfy some criteria of code coverage (e.g., the test designer can create tests to cause all statements in the program to be executed at least once)

Types of testing systems (CI/CD)

- Unit vs. Integration testing
- System testing
- Compatibility testing
- Installation testing
- Smoke and sanity testing
- Regression testing

We will use pytest and fastapi testing system

- <https://fastapi.tiangolo.com/tutorial/testing/>
- <https://docs.pytest.org/>

- 1 4 modules on AWS course (S3, EC2, EBS, RDS) - if you finish all the course you get +10 bonus points to final grade
- 2 Build full REST/HTTP fastapi backend + Dockerization (due 1/4)
- 3 UI (react/streamlit) (due 1/5)
- 4 Docker compose the server with UI and backend plus server and write a clear README with git submodules (due 29/5)
- 5 Presentation of the system in a demo in a 2-3 minutes video on youtube and clear README (due 29/5)


Ideas for projects next semester (based on skills we will learn this semester)

- 1 AI/ML based predictive system
- 2 Smart contracts (web3)
- 3 Any other system with at least 3 microservices

- Monolithic vs. Microservices
- Docker
- Client-Server
- REST/HTTP API
- FastAPI
- Pytest
- asyncio
- Frontend (React javascript and Streamlit python)
- Docker compose
- Functional programming
- How to compile a new library

What Are The Best Software Engineering Principles?

Measure twice and cut once



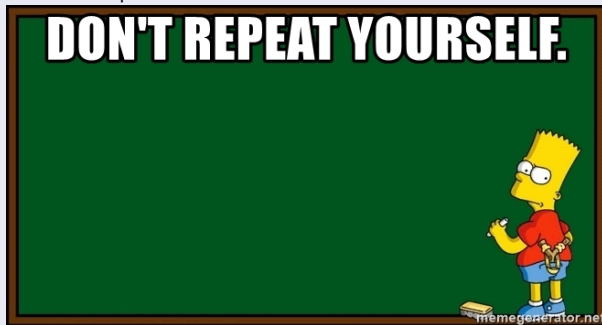
Based on this nice post

<https://luminousmen.com/post/what-are-the-best-engineering-principles>

What Are The Best Software Engineering Principles?

Don't Repeat Yourself (DRY)

If any code occurs more than twice in the codebase, you should think of moving it in a separate function. In fact, you should consider creating a separate method even if you encounter repetition a second time.



What Are The Best Software Engineering Principles?

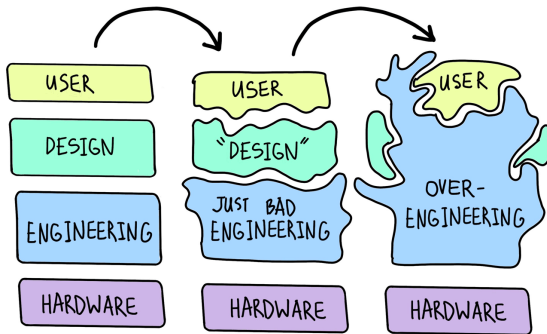
Keep It Simple -Stupid- (KISS)

Some think that this idea transformed from Occam's Razor philosophical principle. You can interpret it as follows: one should not create extra entities to the system without a strong necessity. It is always a good idea to first consider the usefulness of adding another method/class/tool/process, etc.

What Are The Best Software Engineering Principles?

You Aren't Gonna Need It (YAGNI)

Don't implement all the "necessary" (most likely unnecessary) functionality at once from the very beginning of the project.



@luminousmen.com

What Are The Best Software Engineering Principles?

Avoid Premature Optimization

“Premature optimization is the root of all evil (or at least most of it) in programming” — Donald Knuth

Watch Knuth on a talk with Lex Friedman <https://www.youtube.com/watch?v=EE1R8FYUJm0>

Principle Of Least Astonishment

This principle means that your code should be intuitive and obvious, and not surprise another developer when reviewing the code.

Law of Demeter (Olympian goddess of the harvest and agriculture)

The basic idea here is to divide the areas of responsibility between classes and encapsulate the logic within a class, method, or structure.

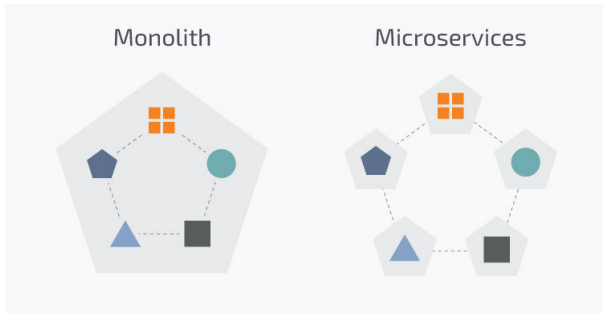
- ❶ **Decoupling** You should try to reduce the number of connections between different classes or entities
- ❷ **Cohesion** The associated classes must be in one module/package/directory

SOLID - create code that is easy to maintain and extend over time

- Single responsibility states that every module or class should have responsibility for a single part of the functionality and that responsibility should be entirely encapsulated by the class
- Open-closed states that software entities (classes, modules, functions, etc.) should be open for extension, but closed for modification
- Liskov substitution states that any inherited class should complement (substitutable), not replace, the behavior of the base class
- Interface segregation states that no client of the class should be forced to depend on methods it does not use
- Dependency inversion says that programmers should work at the interface level and not at the implementation level

Monolithic vs. Microservices

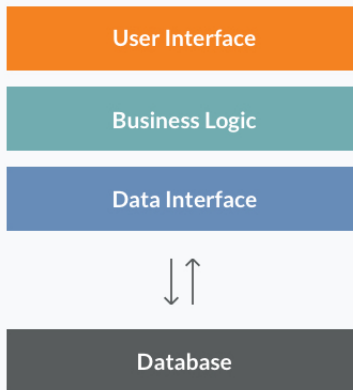
- Monolithic application is a single unified unit that contains all the logic in one entity
- Microservice architecture breaks the application down into a collection of smaller independent units



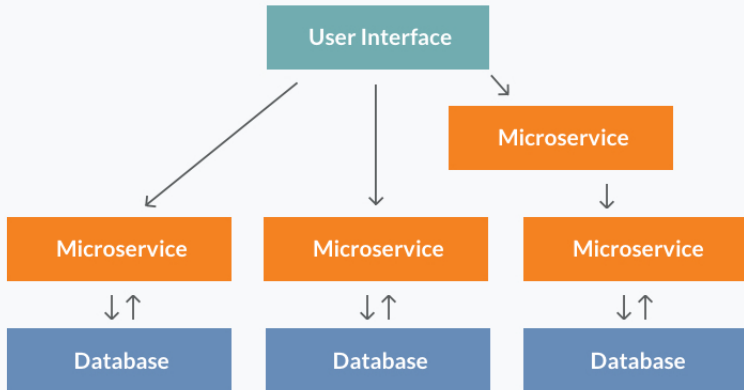
Further reading material

<https://www.n-ix.com/microservices-vs-monolith-which-architecture-best-choice-your-business>

Monolithic Architecture



Microservice Architecture



- Dockerhub/Registry
- Dockerfile
- `docker build`
- `docker run`
- `docker ps`
- `docker network ls`
- `docker volumes`
- `docker expose ports`
- `docker images`
- `docker exec`
- `docker image prune -a`

Further training material

<https://training.play-with-docker.com/alacart/>

<https://towardsdatascience.com/twenty-one-techniques-and-five-concepts-for-better-docker-usage-9ee135dccc9>

Advanced “Setting up a reverse proxy server” (How Docker makes our life easier)

“A very common scenario for developers, is to run their server behind a reverse proxy that sits in front of web servers and forwards client/frontend (e.g., web browser) requests to the web servers (“backend”). There are many reasons why you would want to do this but one of the main reasons is to run your API server on a different network or IP then your front-end application is on. You can then secure this network and only allow traffic from the reverse proxy server. For the sake of simplicity and space, I've created a simple frontend application in React.js and a simple backend API written in Node.js. Run the following command to pull the code from GitHub.”

Without docker this is not easy to do.

Soon we will see how Docker simplifies the process of building a server that run a reverse proxy

Go through (reverse proxy, react, nginx):

<https://www.docker.com/blog/how-to-use-the-official-nginx-docker-image/>

<https://www.docker.com/blog/how-to-use-the-official-nginx-docker-image/>

- ❶ `docker run -it --rm -d -p 8080:80 --name web nginx`
- ❷ `curl http://localhost:8080`
- ❸ `docker stop web`
- ❹ Add `index.html` to local `site-content` and map it to `/usr/share/nginx/html`
(<https://gist.github.com/chrisvfritz/bc010e6ed25b802da7eb>)
- ❺ `docker run -it --rm -d -p 8080:80 --name web -v
~/site-content:/usr/share/nginx/html nginx`
- ❻ Doing stuff via Dockerfile (`docker build -t webserver`):
`FROM nginx:latest`
`COPY ./index.html /usr/share/nginx/html/index.html`
- ❼ `docker run -it --rm -d -p 8080:80 --name web webserver`

Recap - linux commands everyone should know

host/network commands

- ip
- ifconfig
- hostname
- whoami
- uname
- ping

file-related commands

- mkdir rmdir cp mv, rm
- cd ls -l
- find
- wc
- xxd
- du -h /
- chown chmod

archives:

- zip, tar gzip, unzip, gunzip
- tar -czvf name-of-archive.tar.gz /path/to/dir-or-file
c - create, v verbose, f - allow to chose the name

more commands

- cat touch echo
- locate whereis which find
- grep
- df du
- awk head tail
- diff
- jobs (to see background jobs command&)
- kill (sending signals to processes)

SIGTERM (15) – requests the job to stop

SIGKILL (9) – forces programs to stop

- wget curl
- top, htop, brew (should be installed), apt-install

How microservices talk to each other

What is API?

API stands for Application Programming Interface. This interface allows users to build upon another application's functionality.

What is web API?

Web API is when other SW services use other application's/service's functionality over the web/network.

What is HTTP?

HTTP stands for Hypertext Transfer Protocol: an application layer protocol in the Internet protocol suite model for distributed, collaborative, hypermedia information systems. - <http://facebook.com> - <https://facebook.com>

How microservices talk to each other

GET method

```
GET /microservice/v1/function?param1=value1&param=value2
```

POST method

```
POST /microservice/v1/function HTTP/1.1
```

```
Host: localhost
```

```
param1=value1&param=value2
```

How do we perform HTTP requests (postman and cli)

- 1 Postman <https://web.postman.co/home>
- 2 curl or wget in the command line

How do we perform HTTP requests (python)

③ requests library in python:

```
>> r = requests.get('https://api.github.com/user', auth=('user', 'pass'))
>> r.status_code
200
>> r.headers['content-type']
'application/json; charset=utf8'
>> r.encoding
'utf-8'
>> r.text
'{"type":"User"...}'
>> r.json()
{'private_gists': 419, 'total_private_repos': 77, ...}
```

④ httpx library in python <https://www.python-httpx.org/quickstart/>

- 1 Perform GET `https://httpbin.org/get`
- 2 Perform POST `https://httpbin.org/post` with `data={'key': 'value'}`
 - Note that POST/GET could be “overloaded” (have the same endpoint)

What is a REST API?

“When a client request is made via a RESTful API, it transfers a representation of the state of the resource to the requester or endpoint. This information, or representation, is delivered in one of several formats via HTTP: JSON (Javascript Object Notation), HTML, XML, Python, PHP, or plain text. JSON is the most generally popular file format to use because, despite its name, it’s language-agnostic, as well as readable by both humans and machines”

Read more about REST and HTTP

<https://www.redhat.com/en/topics/api/what-is-a-rest-api>
<https://www.educative.io/blog/what-are-rest-apis>

Please complete due next class (March 7th, 2022) and use Discord for help

- 1 Create a remote git repo on our organization GitHub <https://github.com/EASS-HIT-2022/> (private/public)
- 2 Name the repo `http-api-demo-<your github name>`
- 3 Include a README, Dockerfile, client.py files
- 4 In client.py include at least two POST/GET requests from httpbin demo HTTP API (<http://httpbin.org/>):
 - POST to any endpoint of your choice (e.g., <http://httpbin.org/post>)
 - GET to any endpoint of your choice (e.g., <http://httpbin.org/get>)
- 5 Make a Dockerfile that execute client.py on startup and prints the status and output from the http requests it performs from step 3. Helpful snippet:

```
FROM ubuntu
RUN apt-get update
RUN apt-get -y install python
CMD ["echo", "Hello, EASS 2022"]
```

- 6 Now, create a second Dockerfile in your git repo `localhost.Dockerfile` which call to the local hosted httpbin and call it via `http://localhost: docker run -p 80:80 kennethreitz/httpbin`
- 7 build the second docker image. Useful command:

```
docker build -t tab ./ -f localhost.Dockerfile
```

Install

- 1 pip install fastapi
- 2 pip install "uvicorn[standard]" Uvicorn is an ASGI (Asynchronous Server Gateway Interface) web server implementation for Python

code of the server (main.py)

```
from fastapi import FastAPI
```

```
app = FastAPI()
```

```
@app.get("/")
```

```
async def get_root():
```

```
    return {"message": "Hello World", "method": "GET"}
```

```
@app.post("/")
```

```
async def post_root():
```

```
    return {"message": "Hello World", "method": "POST"}
```

Running the server

```
uvicorn main:app --reload
```

read more

<https://fastapi.tiangolo.com/tutorial/>

- 1 Build only the backend (using FastAPI)
- 2 Include a Dockerfile, README and the source code of the app
- 3 Be OOP-friendly (recall the SOLID principle) and use pydantic
- 4 Include both integration and unit tests inside using pytest, httpx, or pip install docker (you may use some bash scripting as well). The idea is to be robust, simple and test the whole system wisely and efficiently.

Suggested layout of the repo:

```
.  
|- app  
|  
|   |- main.py  
|   |- unit_tests.py  
|   |- requirements.txt  
|  
|- integration_test.py  
|- Dockerfile  
|- README.md
```

List of ideas ideas for projects

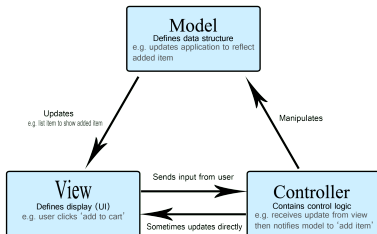
- Building the backend of a voting app
- Personal wallet (keep expenses, images)
- Weather application
- The backend of a US/IL stock viewer analyzer webapp
- The backend of a twitter summarizer webapp (focus on one field e.g., stock symbols)

We want to send data (JSONs) over the wire, but in the code we would like to work with objects.

- DTO (Data Transfer Object)
- ORM (Object Relational Mapping)
- MVC (Model View Controller)

The MVC design pattern (mostly related to UI)

- Model The model manages the data, logic and rules of the application.
- View Any representation of information such as a chart, diagram or table.
- Controller Accepts input and converts it to commands for the model or view.



ref (Mozilla is great for web resources) <https://developer.mozilla.org/en-US/docs/Glossary/MVC/model-view-controller-light-blue.png>

pydantic helps us to define what type of JSONs are valid and what are their interperation. With Pydantic's Model classes we can define the input/outputs of each API endpoint. It helps fastapi with validation, serialization, and documentation.

- ❶ Between microservices (request vs. response)
- ❷ Between Databases microservices
- ❸ Or between any two entities or the user and the application

```
from pydantic import BaseModel
```

```
class User(BaseModel):  
    id: int  
    name = 'Jane Doe'
```

```
from pydantic import BaseModel
```

```
class Client(BaseModel):  
    id: int  
    balance: float
```

```
class Transaction(BaseModel):  
    from_client: Client  
    to_client: Client  
    amount: float
```

```
class Request(BaseModel):  
    id: int  
    transaction: Transaction
```

```
class Response(Request):  
    approved: bool  
    executed: bool
```

```
@app.post("/v1/handle")
def handle(req: Request):
    if req.from_client.balance > req.transaction.amount:
        pass # do something

    res = Response()

    res.id = req.id
    return res
```

Using dataclasses, understanding what's pydantic is doing

```
from pydantic.dataclasses import dataclass
import json
```

```
@dataclass
class User:
    id: int
    name: str
```

```
user = User(id=123, name="James")
d = asdict(user) # {'id': 123, 'name': 'James'}
user_json = json.dumps(d)
print(user_json) # '{"id": 123, "name": "James"}'
```

```
# Or directly with pydantic_encoder
json.dumps(user, default=pydantic_encoder)
```

```
json_raw = '{"id": 123, "name": "James"}'
user_dict = json.loads(json_raw)
user = User(**user_dict)
```

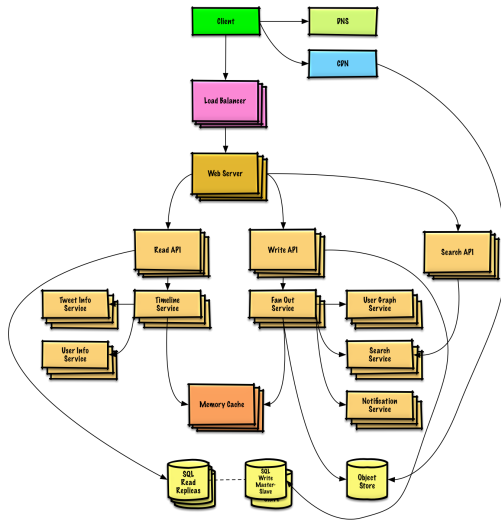
```
user = User.__pydantic_model__.parse_raw('{"id": 123, "name": "James"}')
print(user)
```

ref:

<https://stackoverflow.com/questions/67621046/initializing-a-pydantic-dataclass-from-json>

How to design a project ()

- architecture
- draw.io



Step 1: Outline use cases, constraints, and assumptions

Gather requirements and scope the problem. Ask questions to clarify use cases and constraints. Discuss assumptions.

- Who is going to use it?
- How are they going to use it?
- How many users are there?
- What does the system do?
- What are the inputs and outputs of the system?
- How much data do we expect to handle?
- How many requests per second do we expect?
- What is the expected read to write ratio?

ref

<https://github.com/donnemartin/system-design-primer>

Step 2: Create a high level design

- Outline a high level design with all important components.
- Sketch the main components and connections
- Justify your ideas

How to use draw.io: <https://reneelin2019.medium.com/drawing-cloud-architectures-neural-network-diagrams-and-more-with-draw-io-4f7128ee1aea>

Step 3: Design core components

Dive into details for each core component. For example, if you were asked to design a url shortening service, discuss:

- Generating and storing a hash of the full url
 - MD5 and Base62
 - Hash collisions
 - SQL or NoSQL
 - Database schema
- Translating a hashed url to the full url
 - Database lookup
- API and object-oriented design between the microservices

ref

<https://github.com/donnemartin/system-design-primer>

Identify and address bottlenecks, given the constraints. For example, do you need the following to address scalability issues?

- Load balancer
- Horizontal scaling
- Caching
- Database sharding
- Blue-green deployment to reduce downtime and risk
- Discuss potential solutions and trade-offs. Everything is a trade-off. Address bottlenecks using principles of scalable system design.

- Use back of the envelope calculations
- Powers of two table

Power	Exact Value	Approx Value	Bytes
7	128		
8	256		
10	1024	1 thousand	1 KB
16	65,536		64 KB
20	1,048,576	1 million	1 MB
30	1,073,741,824	1 billion	1 GB
32	4,294,967,296		4 GB
40	1,099,511,627,776	1 trillion	1 TB

- Latency numbers every programmer should know

Latency Comparison Numbers

L1 cache reference	0.5 ns	
Branch mispredict	5 ns	
L2 cache reference	7 ns	14x L1 cache
Mutex lock/unlock	25 ns	
Main memory reference	100 ns	20x L2 cache, 200x L1 cache
Compress 1K bytes with Zip	10 us	
Send 1 KB bytes over 1 Gbps network	10 us	
Read 4 KB randomly from SSD*	150 us	~1GB/sec SSD
Read 1 MB sequentially from memory	250 us	
Round trip within same datacenter	500 us	
Read 1 MB sequentially from SSD*	1 ms	~1GB/sec SSD, 4X memory
HDD seek	10 ms	20x datacenter roundtrip
Read 1 MB sequentially from 1 Gbps	10 ms	40x memory, 10X SSD
Read 1 MB sequentially from HDD	30 ms	120x memory, 30X SSD
Send packet CA->Israel->CA	200 ms	

ref

<https://github.com/donnemartin/system-design-primer>

How would you implement google search engine (discussion)

<https://softwareengineering.stackexchange.com/questions/38324/how-would-you-implement-google-search>

Testing

- SW testing is the process of evaluating and verifying that a software product or application does what it is supposed to do.
- Integration testing
- Unit testing

Profiling

- **Flat profiler** - computes the average call times (callers & callees)
- **Call-graph profiler** - shows the call times, frequencies of the functions and their call graph

ref

<https://docs.python.org/3/library/profile.html>

<https://stackoverflow.com/questions/582336/how-can-you-profile-a-python-script>

<https://medium.com/@alaminopu.me/profiling-your-python-3-code-8c3f695e62da>

<https://pypi.org/project/pytest-benchmark/>

<https://www.ibm.com/topics/software-testing>

Unit tests are written and run by software developers to ensure that a section of an application (“unit”) meets its design and behaves as intended.

installing

```
pip install pytest
```

hello world

```
# content of test_sample.py
```

```
def inc(x):
```

```
    return x + 1
```

```
def test_pos():
```

```
    assert inc(3) == 5
```

```
def test_neg():
```

```
    assert inc(-1) == 0
```

ref

```
https://realpython.com/pytest-python-testing/ https://docs.pytest.org/en/7.1.x/
```

profiling (“program profiling”, “software profiling”) is a form of dynamic program analysis that measures, for example, the space (memory) or time complexity of a program, the usage of particular instructions, or the frequency and duration of function calls. Most commonly, profiling information serves to aid program optimization, and more specifically, performance engineering.

```
import cProfile

def fibonnaci(n):
    if n in [0, 1]:
        return n
    else:
        return fibonnaci(n-1) + fibonnaci(n-2)

if __name__ == '__main__':
    pr = cProfile.Profile()
    pr.enable()
    fibonnaci(100)
    pr.disable()
    pr.print_stats()
```

profiling code

```
pip install pytest-benchmark
pip install aspectlib

import time
import pytest
class Foo(object):
    def __init__(self, arg=0.01):
        self.arg = arg

    def run(self):
        self.internal(self.arg)

    def internal(self, duration):
        time.sleep(duration)

@pytest.mark.benchmark(
    group="group-name",
    min_time=0.1,
    max_time=0.5,
    min_rounds=5,
    timer=time.time,
    disable_gc=True,
    warmup=False
)
def test_my_stuff(benchmark):
```

❶ main.py

```
from fastapi import FastAPI
```

```
app = FastAPI()
```

```
@app.get("/")
```

```
async def read_main():
```

```
    return {"msg": "Hello World"}
```

❷ test_main.py

```
from fastapi import FastAPI
```

```
from fastapi.testclient import TestClient
```

```
app = FastAPI()
```

```
@app.get("/")
```

```
async def read_main():
```


```
    return {"msg": "Hello World"}
```

```
client = TestClient(app)
```

How does something which feels concurrent uses a single thread and a single



CPU?



Event loop

ref

<https://tenthousandmeters.com/blog/python-behind-the-scenes-12-how-asyncawait-works-in-python/>

```
import threading
import concurrent.futures
import time
from tqdm.asyncio import trange, tqdm
import asyncio
import numpy as np

def run_threading(n_threads=5):
    threads = []
    print("Starting...")
    start = time.time()
    for i in range(n_threads):
        thread = threading.Thread(target=print, args=[f"I am thread {i}."])
        thread.start()
        threads.append(thread)

    for thread in threads:
        thread.join()
    end = time.time()
    print(f"Time to complete: {end - start}")
```

```
import threading
import concurrent.futures
import time
from tqdm.asyncio import trange, tqdm
import asyncio
import numpy as np

def do_concurrent(n=32):
    start = 10_000_000
    print_list = [i for i in range(start, start + n)]
    print("Starting...")
    start = time.time()
    with concurrent.futures.ProcessPoolExecutor(max_workers=2) as executor:
        futures = {executor.submit(print, i): i for i in print_list}
    for f in concurrent.futures.as_completed(futures):
        print(f"done {futures[f]} {f.result()}")
    end = time.time()
    print(f"Time to complete: {end - start}")
```

Multithreading, Processes, Asyncio

```
import threading
import concurrent.futures
import time
from tqdm.asyncio import trange, tqdm
import asyncio
import numpy as np

async def wait_and_print(t, n):
    await asyncio.sleep(t)
    print(f"coroutine {n} slept for {t} seconds")

async def do_tqdm_asyncio(n=10):
    arr = []
    async for i in trange(n):
        print(f"asyncio {i}")
        arr.append(wait_and_print(np.random.randint(1, 5), i))
    await asyncio.gather(*arr)

if __name__ == "__main__":
    run_threading(100)
    do_concurrent(100)
    asyncio.run(do_tqdm_asyncio(100))
```


The core behind asyncio are select and poll OS syscalls


```
import selectors
import socket

sel = selectors.DefaultSelector()

def accept(sock, mask):
    conn, addr = sock.accept() # Should be ready
    print('accepted', conn, 'from', addr)
    conn.setblocking(False)
    sel.register(conn, selectors.EVENT_READ, read)

def read(conn, mask):
    data = conn.recv(1000) # Should be ready
    if data:
        print('echoing', repr(data), 'to', conn)
        conn.send(data) # Hope it won't block
    else:
        print('closing', conn)
        sel.unregister(conn)
        conn.close()

sock = socket.socket()
sock.bind(('localhost', 1234))
sock.listen(100)
sock.setblocking(False)
sel.register(sock, selectors.EVENT_READ, accept)
```



The image shows the GitHub Copilot landing page. At the top left is the GitHub Copilot logo. At the top right is a 'Sign up >' button. In the center, there is a 'Technical Preview' badge. Below this is the main heading 'Your AI pair programmer' in large white font. Underneath the heading is the text 'With GitHub Copilot, get suggestions for whole lines or entire functions right inside your editor.' and another 'Sign up >' button. At the bottom, there is a screenshot of a code editor with several tabs: 'sentiment.ts', 'write_sql.go', 'parse_expenses.py', and 'addresses.rb'. The 'sentiment.ts' tab is active, showing a TypeScript code snippet with a comment and an import statement.

GitHub Copilot

Sign up >

Technical Preview

Your AI pair programmer

With GitHub Copilot, get suggestions for whole lines or entire functions right inside your editor.

Sign up >

```
1 #!/usr/bin/env ts-node
2
3 import { fetch } from "fetch-h2";
4
5 // Determine whether the sent
6
7
8
9
10
```



Be active on EASS discord and try to learn and help each other as much as you can.

node.js react <https://docs.microsoft.com/en-us/visualstudio/docker/tutorials/docker-tutorial>
<https://github.com/docker/awesome-compose/tree/master/fastapi>
<https://luminousmen.com/post/what-are-the-best-engineering-principles>
<https://developer.mozilla.org/en-US/docs/Learn/JavaScript/Asynchronous/Concepts> gdb
profiling debugging ipdb basic security <https://owasp.org/www-project-top-ten/>
<https://owasp.org/Top10/>